

Facilitating Kernel UAF Vulnerability Exploitation: an idea

韩天硕

Table of Content

- **Overview**
- **Example & Enlights**
- **Possible Solution**

Overview

UAF vulnerability pattern:

- UAF chunk size
- certain bytes readable
- certain bytes writable

Overview

Given an UAF vul pattern, I want to:

- **identify candidate objects that can be sprayed to the UAF region**
- **figure out syscalls with proper arguments to perform the spray**
- **find proper approach to exploit**

Example1

OCTF 2021 Final - kernote

UAF vulnerability pattern:

- UAF size: one kmalloc-32 chunk
- readable bytes: none
- writable bytes: the first 8 bytes

exploitation: *ldt_struct*

Example1

OCTF 2021 Final - kernote

```
struct ldt_struct {
    /*
     * Xen requires page-aligned LDTs with special permissions. This is
     * needed to prevent us from installing evil descriptors such as
     * call gates. On native, we could merge the ldt_struct and LDT
     * allocations, but it's not worth trying to optimize.
     */
    struct desc_struct    *entries;
    unsigned int          nr_entries;

    /*
     * If PTI is in use, then the entries array is not mapped while we're
     * in user mode. The whole array will be aliased at the addressed
     * given by ldt_slot_va(slot). We use two slots so that we can allocate
     * and map, and enable a new LDT without invalidating the mapping
     * of an older, still-in-use LDT.
     */
    int                    slot;
};
```

```
SYSCALL_DEFINE3(modify_ldt, int , func , void __user * , ptr ,
                unsigned long , bytecount)
{
    int ret = -ENOSYS;

    switch (func) {
    case 0:
        ret = read_ldt(ptr, bytecount);
        break;
    case 1:
        ret = write_ldt(ptr, bytecount, 1);
        break;
    case 2:
        ret = read_default_ldt(ptr, bytecount);
        break;
    case 0x11:
        ret = write_ldt(ptr, bytecount, 0);
        break;
    }
    /*
     * The SYSCALL_DEFINE() macros give us an 'unsigned long'
     * return type, but the ABI for sys_modify_ldt() expects
     * 'int'. This cast gives us an int-sized value in %rax
     * for the return code. The 'unsigned' is necessary so
     * the compiler does not try to sign-extend the negative
     * return codes into the high half of the register when
     * taking the value from int->long.
     */
    return (unsigned int)ret;
}
```

Example1

OCTF 2021 Final - kernote

```
static int read_ldt(void __user *ptr, unsigned long bytecount)
{
    struct mm_struct *mm = current->mm;
    unsigned long entries_size;
    int retval;

    .....

    if (copy_to_user(ptr, mm->context.ldt->entries, entries_size)) {
        retval = -EFAULT;
        goto out_unlock;
    }

    .....
```

Example1

OCTF 2021 Final - kernote

two solves during the contest, both unintended:

- *@Balsn* gave a 1-day exp, unrelated to the vulnerability at all
- *@Organizor* gave a kernel ROP solve, too delicate and complicated to reproduce

conclusion: Identifying proper structs is vital for UAF exploitation

Example1

OCTF 2021 Final - kernote

bitbucket.org

なお、検証にはLinux Kernel 4.19.98を使いました。

- [はじめに](#)
- [Leak/AAR/AW/RIP制御に使える構造体一覧](#)
 - [shm_file data](#)
 - [seq_operations](#)
 - [msg_msg_\(+user-supplied data\)](#)
 - [subprocess_info](#)
 - [cred](#)
 - [file](#)
 - [timerfd_ctx](#)
 - [tty_struct](#)
- [任意データ書き込み/Heap_Sprayに使える構造体](#)
 - [msg_msg](#)
 - [setxattr](#)
 - [sendmsg](#)
- [参考文献](#)

Q1: where do these structs come from?

Example2

CVE-2021-26708

UAF vulnerability pattern:

- **UAF size: one kmalloc-64 chunk**
- **readable bytes: none**
- **writable bytes: 4 bytes at offset 40**

Example2

CVE-2021-26708

```
/* one msg_msg structure for each message */
struct msg_msg {
    struct list_head m_list;
    long m_type;
    size_t m_ts;           /* message text size */
    struct msg_msgseg *next;
    void *security;
    /* the actual message follows immediately */
};
```

the pointer *security* is at offset 40, and it will be passed to kfree() resulting to arbitrary free, and can be transformed into UAF

Example2

CVE-2021-26708

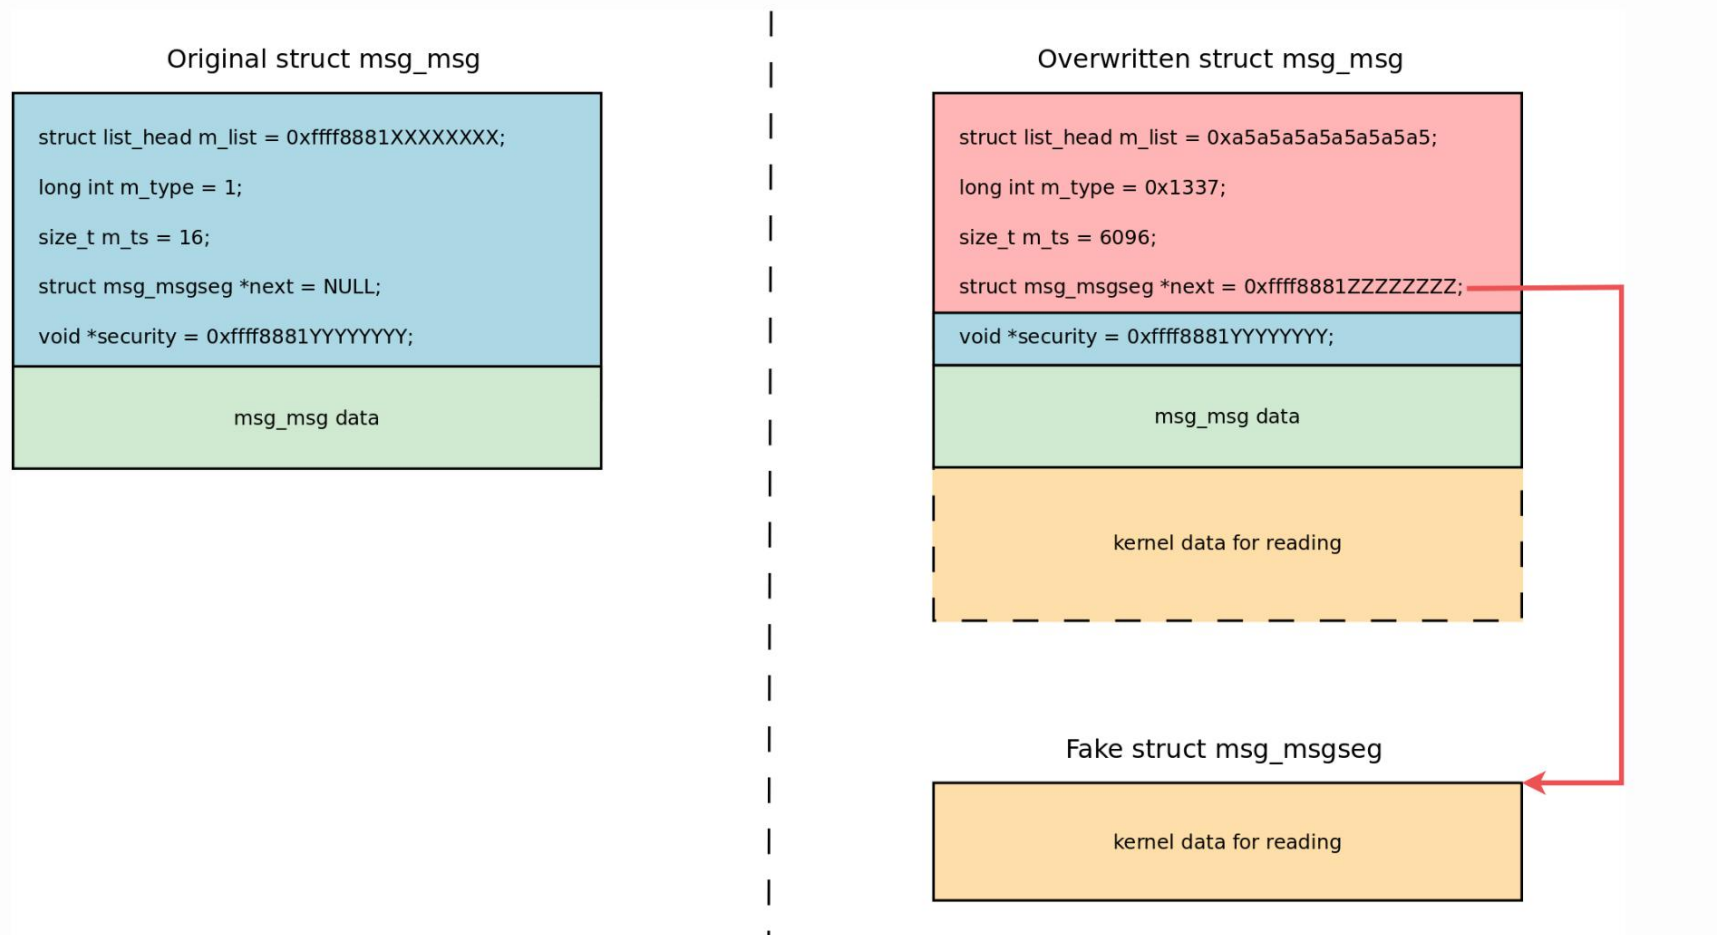
New UAF vulnerability pattern:

- **UAF size: one kmalloc-64 chunk**
- **readable bytes: none**
- **writable bytes: the first 40 bytes**

Example2

CVE-2021-26708

arbitrary read:



Example2

CVE-2021-26708

execution path around msg_msg struct can be extended, achieving both **arbitrary read and write** (corCTF 2021 Finals - Fire of Salvation)

Q2: finding such execution path needs extensive expertise in kernel and intensive manual efforts. **Can it be automatized?**

Example3

QWB2021 - notebook

UAF vulnerability pattern:

- **UAF size: one kmalloc-0x400 chunk**
- **readable bytes: the first 0x60 bytes**
- **writable bytes: the first 0x60 bytes**

Example3

QWB2021 - notebook

a fancy function, available in kernels who enable multi-cpu support:

```
static void work_for_cpu_fn(struct work_struct *work)
{
    struct work_for_cpu *wfc = container_of(work, struct work_for_cpu, work);

    wfc->ret = wfc->fn(wfc->arg);
}
```

after compilation:

```
static void work_for_cpu_fn(size_t * args)
{
    args[6] = ((size_t (*) (size_t)) (args[4])(args[5]));
}
```


Example3

QWB2021 - notebook

perfect for us to execute *commit_creds(prepare_kernel_cred(0))*

- hijack *ioctl()* of a tty struct to *work_for_cpu_fn()*
- set *tty_struct[4]=prepare_kernel_cred*, *tty_struct[5]=0*
- invoke *ioctl()* and the new generated cred struct will be placed at *tty_struct[6]*, and we can read it
- do the trick again to invoke *commit_creds()*

Example3

QWB2021 - notebook

conclusion: execution path of exploit is not limited to syscall sequences, sometimes can be delicate.

Q3: Is it possible to identify more such fancy functions?

Possible Solution

Q1: where do these structs come from?

- manually create an UAF chunk
- trace kmalloc()
- fuzz the kernel until the UAF chunk is taken

Possible Solution

Q2: finding such execution path needs extensive expertise in kernel and intensive manual efforts. **Can it be automatized?**

one possible solution...

if we can identify all control flow related to the struct, we can symbolically execute them until we meet instructions like

- `call/jmp SYMBOLIC_VALUE` ----- **control flow hijack**
- `mov [SYMBOLIC_VALUE], rsi` ----- **arbitrary read**
-

Possible Solution

Q3: Is it possible to identify more such fancy functions?

heuristic rules

Related work

- **Wu, W., Chen, Y., Xu, J., Xing, X., Gong, X., & Zou, W. (2018). FUZE: Towards facilitating exploit generation for kernel use-after-free vulnerabilities. Proceedings of the 27th USENIX Security Symposium, 781–797.**
- **Chen, Y., & Xing, X. (2019). Slake: Facilitating slab manipulation for exploiting vulnerabilities in the linux kernel. Proceedings of the ACM Conference on Computer and Communications Security, 1707–1722. <https://doi.org/10.1145/3319535.3363212>**

Thanks!