

# Facilitating Linux Kernel Heap Vulnerability Exploitation: an idea

韩天硕

# **Table of Content**

- **Overview**
- **Example & Enlights**
- **Related work**
- **unsolved challenges & thingking**
- **Possible Solution**

# Overview

## Kernel heap vulnerability form:

- double free
- use after free
- out of bound read / write
- .....

# Overview

**Kernel heap vulnerability capability model:**

- the chunk size
- readable/writable bytes within(around) the chunk

# Overview

**General pattern to exploit a heap vul model:**

- **identify an useful kernel object from kernel code space**
- **allocate such an object to region under control**
- **utilizing an execution path(related to the object) to finish exploitation / help with exploitation**

# Overview

**General patterns to exploit a heap vul model:**

- **identify an useful kernel object from kernel code space**
- **allocate such an object to region under control**
- **utilizing an execution path(related to the object) to finish exploitation / help with exploitation**

**all are unsolved challenges**

# Overview

**I want to build a database/tool that, given a vul model, can indicate me: what can be done with such a model**

**more specifically:**

- list candidate objects, that can fit in the evil chunk size and may be useful for exploitation**
- syscall sequences to allocate such useful objects**
- list the objects' capability and syscall sequences to trigger**

# Example1

## OCTF 2021 Final - kernote

```
void* kernote_ioctl(void* d, long cmd, long arg){
    ....
    void* ptr;
    void* note[0x10];

    if(cmd == CMD_ADD){
        note[arg] = kmalloc(0x20);
    }else if(cmd == CMD_SELECT){
        ptr = note[arg];
    }else if(cmd == CMD_WRITE){
        copy_from_user(ptr, arg, 8);
    }else if(cmd == CMD_DELE){
        kfree(note[arg]);
    }
    ....
}
```



# Example1

OCTF 2021 Final - kernote

**vulnerability model:**

- **chunk size: one kmalloc-32 chunk**
- **readable bytes: none**
- **writable bytes: the first 8 bytes**

**exploitation: *ldt\_struct***

# Example1

## OCTF 2021 Final - kernote

```
struct ldt_struct {
    /*
     * Xen requires page-aligned LDTs with special permissions. This is
     * needed to prevent us from installing evil descriptors such as
     * call gates. On native, we could merge the ldt_struct and LDT
     * allocations, but it's not worth trying to optimize.
     */
    struct desc_struct    *entries;
    unsigned int          nr_entries;

    /*
     * If PTI is in use, then the entries array is not mapped while we're
     * in user mode. The whole array will be aliased at the addressed
     * given by ldt_slot_va(slot). We use two slots so that we can allocate
     * and map, and enable a new LDT without invalidating the mapping
     * of an older, still-in-use LDT.
     */
    int                    slot;
};
```

```
SYSCALL_DEFINE3(modify_ldt, int , func , void __user * , ptr ,
                unsigned long , bytecount)
{
    int ret = -ENOSYS;

    switch (func) {
    case 0:
        ret = read_ldt(ptr, bytecount);
        break;
    case 1:
        ret = write_ldt(ptr, bytecount, 1);
        break;
    case 2:
        ret = read_default_ldt(ptr, bytecount);
        break;
    case 0x11:
        ret = write_ldt(ptr, bytecount, 0);
        break;
    }
    /*
     * The SYSCALL_DEFINE() macros give us an 'unsigned long'
     * return type, but the ABI for sys_modify_ldt() expects
     * 'int'. This cast gives us an int-sized value in %rax
     * for the return code. The 'unsigned' is necessary so
     * the compiler does not try to sign-extend the negative
     * return codes into the high half of the register when
     * taking the value from int->long.
     */
    return (unsigned int)ret;
}
```

# Example1

OCTF 2021 Final - kernote

```
static int read_ldt(void __user *ptr, unsigned long bytecount)
{
    struct mm_struct *mm = current->mm;
    unsigned long entries_size;
    int retval;

    .....

    if (copy_to_user(ptr, mm->context.ldt->entries, entries_size)) {
        retval = -EFAULT;
        goto out_unlock;
    }

    .....
```

# Example1

OCTF 2021 Final - kernote

two solves during the contest, both unintended:

- *@Balsn* gave a 1-day exp, unrelated to the vulnerability at all
- *@Organizor* gave a kernel ROP solve, too delicate and complicated to reproduce

conclusion: Identifying proper structs and its useful execution path is vital for kernel heap exploitation

# Example1

OCTF 2021 Final - kernote

[bitbucket.org](https://bitbucket.org)

なお、検証にはLinux Kernel 4.19.98を使いました。

- [はじめに](#)
- [Leak/AAR/AW/RIP制御に使える構造体一覧](#)
  - [shm\\_file data](#)
  - [seq\\_operations](#)
  - [msg\\_msg\\_\(+user-supplied data\)](#)
  - [subprocess\\_info](#)
  - [cred](#)
  - [file](#)
  - [timerfd\\_ctx](#)
  - [tty\\_struct](#)
- [任意データ書き込み/Heap\\_Sprayに使える構造体](#)
  - [msg\\_msg](#)
  - [setxattr](#)
  - [sendmsg](#)
- [参考文献](#)

**Q1: where do these structs and execution path come from?**

# **Example2**

CVE-2021-26708

**vulnerability model:**

- **chunk size: one kmalloc-64 chunk**
- **readable bytes: none**
- **writable bytes: 4 bytes at offset 40**

# Example2

CVE-2021-26708

```
/* one msg_msg structure for each message */
struct msg_msg {
    struct list_head m_list;
    long m_type;
    size_t m_ts;           /* message text size */
    struct msg_msgseg *next;
    void *security;
    /* the actual message follows immediately */
};
```

the pointer *security* is at offset 40, and it will be passed to kfree() resulting to arbitrary free, and can be transformed into UAF

# **Example2**

**CVE-2021-26708**

**Combining with the vul other capabilities, the vulnerability pattern can be transformed:**

- **UAF size: one kmalloc-64 chunk**
- **readable bytes: none**
- **writable bytes: the first 40 bytes**



# **Example2**

CVE-2021-26708

**Conclusion: some execution path does not follow general exploitation paradigm, but may work at specific vul context**

**Q2: finding such execution path needs extensive expertise in kernel and intensive manual efforts. Can it be automatized?**

# **Example3**

QWB2021 - notebook

**vulnerability model:**

- **UAF size: one kmalloc-0x400 chunk**
- **readable bytes: the first 0x60 bytes**
- **writable bytes: the first 0x60 bytes**

# Example3

QWB2021 - notebook

a fancy function, available in kernels who enable multi-cpu support:

```
static void work_for_cpu_fn(struct work_struct *work)
{
    struct work_for_cpu *wfc = container_of(work, struct work_for_cpu, work);

    wfc->ret = wfc->fn(wfc->arg);
}
```

after compilation:

```
static void work_for_cpu_fn(size_t * args)
{
    args[6] = ((size_t (*) (size_t)) (args[4])(args[5]));
}
```

# Example3

QWB2021 - notebook

perfect for us to execute *commit\_creds(prepare\_kernel\_cred(0))*

- hijack *ioctl()* of a tty struct to *work\_for\_cpu\_fn()*
- set *tty\_struct[4]=prepare\_kernel\_cred*, *tty\_struct[5]=0*
- invoke *ioctl()* and the new generated cred struct will be placed at *tty\_struct[6]*, and we can read it
- do the trick again to invoke *commit\_creds()*

# **Example3**

QWB2021 - notebook

**conclusion 1: execution path of exploit is not limited to syscall sequences, sometimes can be subtle.**

**conclusion 2: overread is as important as overwrite.**

**Q3: by taking extra dimensions of vul capability into consideration, is it possible to identify more exploit paradigm?**

# Related Work

- Wu, W., Chen, Y., Xu, J., Xing, X., Gong, X., & Zou, W. (2018). FUZE: Towards facilitating exploit generation for kernel use-after-free vulnerabilities. Proceedings of the 27th USENIX Security Symposium, 781–797.
- Chen, Y., & Xing, X. (2019). **Slake: Facilitating slab manipulation for exploiting vulnerabilities in the linux kernel**. Proceedings of the ACM Conference on Computer and Communications Security
- Chen, Y., Lin, Z., & Xing, X. (2020). **A Systematic Study of Elastic Objects in Kernel Exploitation**. Proceedings of the ACM Conference on Computer and Communications Security

# Related Work: SLAKE

SLAKE: identify kernel objects that enclose at least one **function pointer**, then filter out objects that may be utilized for **RIP hijacking**

- the object can be allocated from user land
- the function pointer can be called from user land

# Related Work: SLAKE

e.g.

```
struct seq_operations {  
    void * (*start) (struct seq_file *m, loff_t *pos);  
    void (*stop) (struct seq_file *m, void *v);  
    void * (*next) (struct seq_file *m, void *v, loff_t *pos);  
    int (*show) (struct seq_file *m, void *v);  
};
```

```
// allocate  
int fd = open("/proc/self/stat");  
  
// trigger  
read(fd, buf, 1);
```



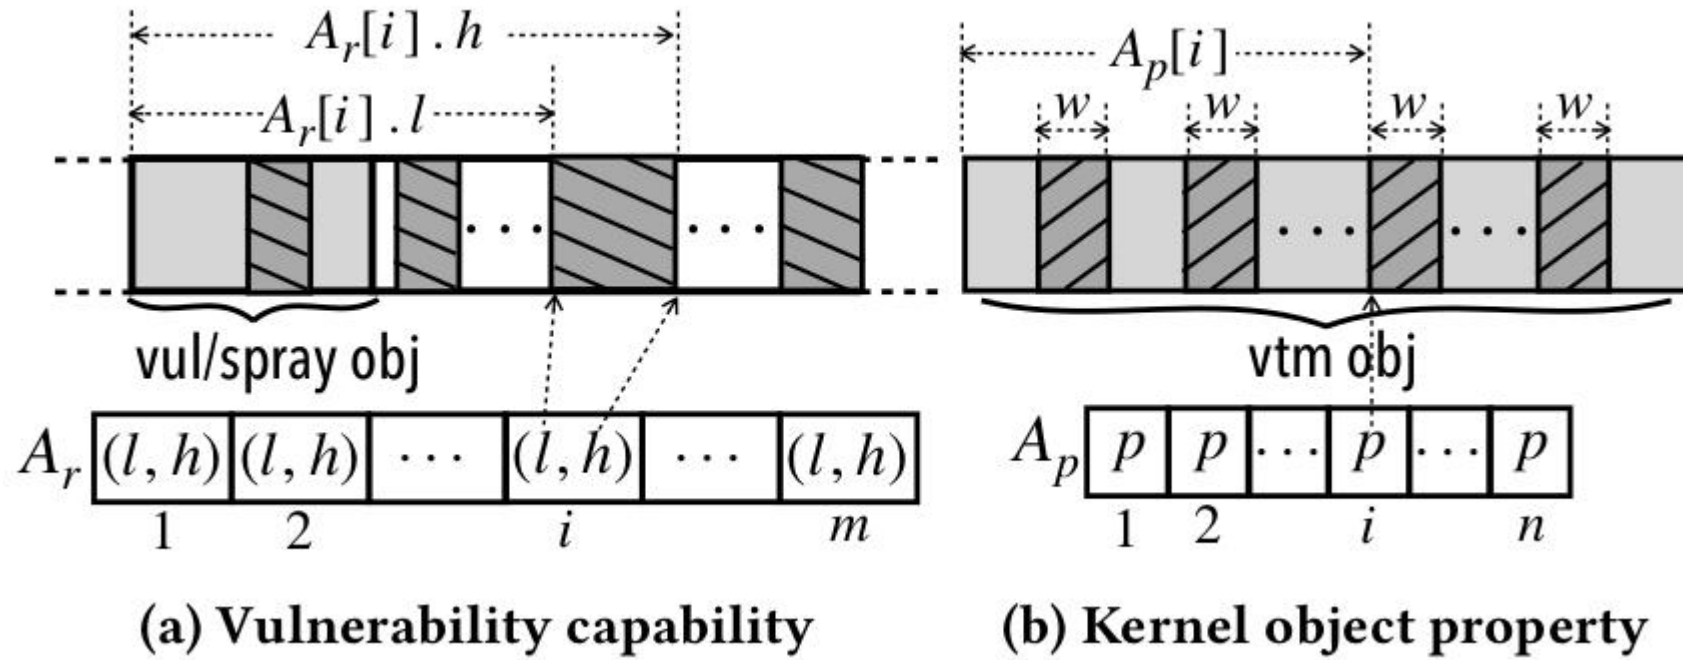
# Related Work: SLAKE

SLAKE answered such a question: I have a heap chunk overwrite primitive and I want to do a **RIP hijacking**, how can I do it?

The identified useful kernel objects are stored in a database, and will be paired with vul capability model

# Related Work: SLAKE

## SLAKE vul capability model



# Related Work: ELOISE

**ELOISE: focuses on elastic kernel objects, identify kernel objects and execution path for **kernel space memory leaking****

**ELOISE answerd such a question: I have a kernel heap chunk overwrite primitive and I want to do a kernel space memory leaking, how can I do it?**

# Related Work: ELOISE

e.g.

```
struct ldt_struct {
    /*
     * Xen requires page-aligned LDTs with special permissions. This is
     * needed to prevent us from installing evil descriptors such as
     * call gates. On native, we could merge the ldt_struct and LDT
     * allocations, but it's not worth trying to optimize.
     */
    struct desc_struct    *entries;
    unsigned int          nr_entries;

    /*
     * If PTI is in use, then the entries array is not mapped while we're
     * in user mode. The whole array will be aliased at the address
     * given by ldt_slot_va(slot). We use two slots so that we can allocate
     * and map, and enable a new LDT without invalidating the mapping
     * of an older, still-in-use LDT.
     *
     * slot will be -1 if this LDT doesn't have an alias mapping.
     */
    int                    slot;
};
```

```
SYSCALL_DEFINE3(modify_ldt, int , func , void __user * , ptr ,
                unsigned long , bytecount)
{
    int ret = -ENOSYS;

    switch (func) {
    case 0:
        ret = read_ldt(ptr, bytecount);
        break;
    case 1:
        ret = write_ldt(ptr, bytecount, 1);
        break;
    case 2:
        ret = read_default_ldt(ptr, bytecount);
        break;
    case 0x11:
        ret = write_ldt(ptr, bytecount, 0);
        break;
    }
    /*
     * The SYSCALL_DEFINE() macros give us an 'unsigned long'
     * return type, but the ABI for sys_modify_ldt() expects
     * 'int'. This cast gives us an int-sized value in %rax
     * for the return code. The 'unsigned' is necessary so
     * the compiler does not try to sign-extend the negative
     * return codes into the high half of the register when
     * taking the value from int->long.
     */
    return (unsigned int)ret;
}
```

# Unsolved challenges & thinking

**1. there is a great number of exploitation paradigms, besides RIP hijacking and kernel space memory leaking**

- reading kernel space memory to a readable region(besides transfer them directly to user space)
- kernel space memory overwrite
- arbitrary kfree()
- leaking not by elastic objects
- critical kernel data corruption
- spraying data into kernel heap chunk
- .....

**digging the kernel according to known exploitation paradigms can be inefficient and unscaleble**

# Unsolved challenges & thinking

## 2. the vul capability model is inadequate to fully describe a vul's capability

- chunk size
  - writable bytes
- chunk size
  - chunk number I can control
  - times I can control for the same chunk
  - is the chunk at same address
  - writable bytes
  - times I can write
  - readable bytes
  - times I can read
  - .....

**stronger the vul capability indicates more possible exploitation path**

# Unsolved challenges & thinking

**3. in many situations, I only know the vul capability at hand, I have no idea what to do next**

**CVE-2021-26708 vulnerability model:**

- chunk size: one kmalloc-64 chunk
- readable bytes: none
- writable bytes: 4 bytes at offset 40

- RIP hijacking is impossible
- kernel memory leaking is unlikely to work
- what I do with this primitive?

**knowing what can be done with a vul primitive can greatly indicate how to perform exploit**

# Possible Solution

**loose the constraint**

**SLAKE and ELOISE digs the kernel strictly subject to constraints, to make sure the found objects and syscall sequences follows the desired exploitation paradigm**



# Possible Solution

mining the kernel not according to **known exploitation paradigms**, but by considering **value** and **reachability**

- **value**: heuristically defining valuable operations
- **reachability**: filter out operations can be triggered by syscall sequences from user land

# Possible Solution

**value:**

1. defining valuable operations: memory read/write, call/jmp, critical kernel functions like *kfree()*, .....
2. mark controllable bytes as taint source / symbolic value
3. checking if valuable operations can be tainted / controlled by symbolic values

# Possible Solution

**reachability:**

- 1. identify all valuable operation sites in execution path related to the kernel object**
- 2. fuzz syscalls, until valuable operation sites reached**

# Possible Solution

e.g. find the arbitrary free() primitive in CVE-2021-26708

```
/* one msg_msg structure for each message */
struct msg_msg {
    struct list_head m_list;
    long m_type;
    size_t m_ts;           /* message text size */
    struct msg_msgseg *next;
    void *security;
    /* the actual message follows immediately */
};
```

1. mark all fields in this struct as taint source
2. fuzz the kernel until found kfree() is tainted by *security*

# Possible Solution

store identified **kernel objects**, **syscalls** and corresponding **capability** in a database, and give an algorithm that retrieve information from the database. Given a vul model such a database can:

- greatly indicate exploitation path by knowing all possible operations
- give indications on how to implement an exploitation path

**Thanks!**